

# Parameterized and Runtime-tunable Snapshot Isolation in Distributed Transactional Key-value Stores

**Hengfeng Wei**, Yu Huang, Jian Lu

Nanjing University, China

September 18, 2017



# Parameterized and Runtime-tunable Snapshot Isolation

## RVSI: Relaxed Version Snapshot Isolation

- 1 Motivation for RVSI
- 2 Definition of RVSI
- 3 CHAMELEON Prototype and RVSI Protocol
- 4 Experimental Evaluation

# Parameterized and Runtime-tunable Snapshot Isolation

## RVSI: Relaxed Version Snapshot Isolation

- 1 Motivation for RVSI
- 2 Definition of RVSI
- 3 CHAMELEON Prototype and RVSI Protocol
- 4 Experimental Evaluation



Figure: Distributed key-value stores.

`put(K key, V val)`      `get(K key)`

# Transactional semantics

existential consistency atomic visibility example

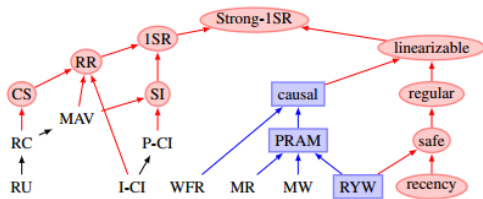


Figure: Transactional consistency models (from [Bailis@VLDB'14]).




Figure: Transactional consistency models (from [Bailis@VLDB'14]).

Snapshot isolation (SI [Berenson@SIGMOD'95], [Adya@Thesis'99]):

- ▶ **Read** from the “latest” snapshot as of the time the transaction started
- ▶ No **write**-conflicting concurrent transactions

Reading the “latest” in a distributed setting  
often requires intensive coordinations.

---

<sup>1</sup>GSI: Generalized Snapshot Isolation [Elnikety@SRDS'05]

<sup>2</sup>NMSI: Non-Monotonic Snapshot Isolation [Ardekani@SRDS'13]

<sup>3</sup>PL-FCV: Forward Consistent View [Aday@Thesis'99]

<sup>4</sup>PSI: Parallel Snapshot Isolation [Sovran@SOSP'11]



Reading the “latest” in a distributed setting often requires intensive coordinations.

Relaxed variants of (distributed) SI:

GSI<sup>1</sup>: allows to read from “older” snapshots

NMSI<sup>2</sup>: allows to observe non-monotonically ordered snapshots

PL-FCV<sup>3</sup>: allows a transaction to observe the updates of transactions that commit after it started

PSI<sup>4</sup>: causal ordering of transactions across sites

---

<sup>1</sup>GSI: Generalized Snapshot Isolation [Elnikety@SRDS'05]

<sup>2</sup>NMSI: Non-Monotonic Snapshot Isolation [Ardekani@SRDS'13]

<sup>3</sup>PL-FCV: Forward Consistent View [Aday@Thesis'99]

<sup>4</sup>PSI: Parallel Snapshot Isolation [Sovran@SOSP'11]

Two possible drawbacks:

1. Unbounded inconsistency

- ▶ no specification of the severity of the anomalies w.r.t SI

Two possible drawbacks:

1. Unbounded inconsistency

- ▶ no specification of the severity of the anomalies w.r.t SI

2. Untunable at runtime

- ▶ determined at the system design phase
- ▶ remain unchanged once the system is deployed

An online bookstore application for motivating  
“bounded inconsistency” and “runtime-tuable”:

Title	Authors	Publisher	Sales	Inventory	Ratings	Reviews	...
-------	---------	-----------	-------	-----------	---------	---------	-----

An online bookstore application for motivating  
 “bounded inconsistency” and “runtime-tuable”:

Title	Authors	Publisher	Sales	Inventory	Ratings	Reviews	...
-------	---------	-----------	-------	-----------	---------	---------	-----

Customer ( $T_1$ ): Obtaining the basic info. about a book  
 ▶ *out-of-date* reviews

An online bookstore application for motivating  
 “bounded inconsistency” and “runtime-tuable”:

Title	Authors	Publisher	Sales	Inventory	Ratings	Reviews	...
-------	---------	-----------	-------	-----------	---------	---------	-----

Customer ( $T_1$ ): Obtaining the basic info. about a book

- ▶ *out-of-date* reviews

Bookstore Clerk ( $T_2$ ): Checking the inventory of a book

- ▶ inventory updated by concurrent transactions committed *after*  $T_2$  starts

An online bookstore application for motivating  
 “bounded inconsistency” and “runtime-tuable”:

Title	Authors	Publisher	Sales	Inventory	Ratings	Reviews	...
-------	---------	-----------	-------	-----------	---------	---------	-----

Customer ( $T_1$ ): Obtaining the basic info. about a book

- ▶ *out-of-date* reviews

Bookstore Clerk ( $T_2$ ): Checking the inventory of a book

- ▶ inventory updated by concurrent transactions committed *after*  $T_2$  starts

Sales Analyst ( $T_3$ ): Studying sales vs. ratings of a book

- ▶ sales and ratings from *separate snapshots*

The idea of “parameterized and runtime-tunable snapshot isolation”.

► RVSI: Relaxed Version Snapshot Isolation

$k_1$ -BV:  $k_1$ -version bounded *backward* view

$k_2$ -FV:  $k_2$ -version bounded *forward* view

$k_3$ -SV:  $k_3$ -version bounded *snapshot* view

---

<sup>5</sup><http://www.aliyun.com/>



The idea of “parameterized and runtime-tunable snapshot isolation”.

- ▶ RVSI: Relaxed Version Snapshot Isolation

  - $k_1$ -BV:  $k_1$ -version bounded *backward* view

  - $k_2$ -FV:  $k_2$ -version bounded *forward* view

  - $k_3$ -SV:  $k_3$ -version bounded *snapshot* view

- ▶ CHAMELEON prototype: distributed transactional key-value store

  - ▶ achieves RVSI

  - ▶ allows each transaction to tune its consistency level at runtime

The idea of “parameterized and runtime-tunable snapshot isolation”.

► RVSI: Relaxed Version Snapshot Isolation

$k_1$ -BV:  $k_1$ -version bounded *backward* view

$k_2$ -FV:  $k_2$ -version bounded *forward* view

$k_3$ -SV:  $k_3$ -version bounded *snapshot* view

► CHAMELEON prototype: distributed transactional key-value store

- achieves RVSI
- allows each transaction to tune its consistency level at runtime
- deployed on Aliyun<sup>5</sup>
- explore the impacts of RVSI on the transaction abort rates

---

<sup>5</sup><http://www.aliyun.com/>

# Parameterized and Runtime-tunable Snapshot Isolation

## RVSI: Relaxed Version Snapshot Isolation

- 1 Motivation for RVSI
- 2 Definition of RVSI
- 3 CHAMELEON Prototype and RVSI Protocol
- 4 Experimental Evaluation

Transaction  $T_i$ :  $s_i \ (r_i/w_i)^+ \ c_i/a_i$

$s_i$ : start operation

$r_i/w_i$ : read/write operation

$c_i/a_i$ : commit/abort operation

Transaction  $T_i$ :  $s_i \ (r_i/w_i)^+ \ c_i/a_i$

$s_i$ : start operation

$r_i/w_i$ : read/write operation

$c_i/a_i$ : commit/abort operation

$x_i$ : version  $i$  of data item  $x$  written by  $T_i$

$r_i(x_j)$ : transaction  $T_i$  reading  $x_j$

$w_i(x_i)$ : transaction  $T_i$  writing  $x_i$

History: modelling an execution of a transactional key-value store

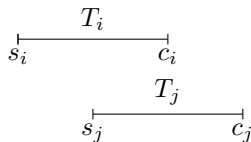
- ▶ *time-precedes partial order*  $\prec_h$  over operations

History: modelling an execution of a transactional key-value store

- ▶ *time-precedes partial order*  $\prec_h$  over operations

Two transactions are *concurrent* if

$$s_i \prec_h c_j \wedge s_j \prec_h c_i$$



A history  $h$  is in *snapshot isolation* iff it satisfies [\[Adya@Thesis'99\]](#)

**Snapshot Read:** All reads of transaction  $T_i$  occur at  $T_i$ 's start time.

**Snapshot Write:** No concurrent committed transactions may write the same data item. (WCF: write-conflict freedom)



A history  $h$  is in *snapshot isolation* iff it satisfies [Adya@Thesis'99]

**Snapshot Read:** All reads of transaction  $T_i$  occur at  $T_i$ 's start time.

$$\begin{aligned} \forall r_i(x_{j \neq i}), w_{k \neq j}(x_k), c_k \in h : \\ (c_j \in h \wedge c_j \prec_h s_i) \wedge (s_i \prec_h c_k \vee c_k \prec_h c_j). \end{aligned}$$

**Snapshot Write:** No concurrent committed transactions may write the same data item. (WCF: write-conflict freedom)

$$\forall w_i(x_i), w_{j \neq i}(x_j) \in h \implies (c_i \prec_h s_j \vee c_j \prec_h s_i).$$

## Principles of RVSI:

- ▶ Using parameters  $(k_1, k_2, k_3)$  to control the severity of the anomalies w.r.t SI

---

<sup>6</sup>RC: Read Committed isolation.

## Principles of RVSI:

- ▶ Using parameters  $(k_1, k_2, k_3)$  to control the severity of the anomalies w.r.t SI
- ▶  $RC^6 \supset RVSI(k_1, k_2, k_3) \supset SI$
- ▶  $RVSI(\infty, \infty, \infty) = RC$        $RVSI(1, 0, *) = SI$

---

<sup>6</sup>RC: Read Committed isolation.

*Each transaction reads data from the “latest” snapshot as of the time the transaction started.*

*– The “Snapshot Read” property of SI*

RVSI relaxes “Snapshot Read” in three ways:

*Each transaction reads data from the “latest” snapshot as of the time the transaction started.*

– *The “Snapshot Read” property of SI*

RVSI relaxes “Snapshot Read” in three ways:

$k_1$ -BV (Backward View): “stale” data versions

staleness  $\leq k_1$

*Each transaction reads data from the “latest” snapshot as of the time the transaction started.*

*– The “Snapshot Read” property of SI*

RVSI relaxes “Snapshot Read” in three ways:

$k_1$ -BV (Backward View): “stale” data versions staleness  $\leq k_1$

$k_2$ -FV (Forward View): “concurrent” data versions forward level  $\leq k_2$

*Each transaction reads data from the “latest” snapshot as of the time the transaction started.*

*– The “Snapshot Read” property of SI*

RVSI relaxes “Snapshot Read” in three ways:

$k_1$ -BV (Backward View): “stale” data versions staleness  $\leq k_1$

$k_2$ -FV (Forward View): “concurrent” data versions forward level  $\leq k_2$

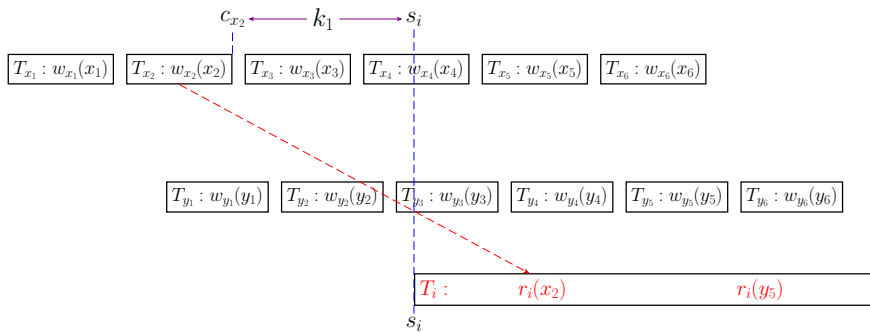
$k_3$ -SV (Snapshot View): “non-snapshot” data versions distance  $\leq k_3$

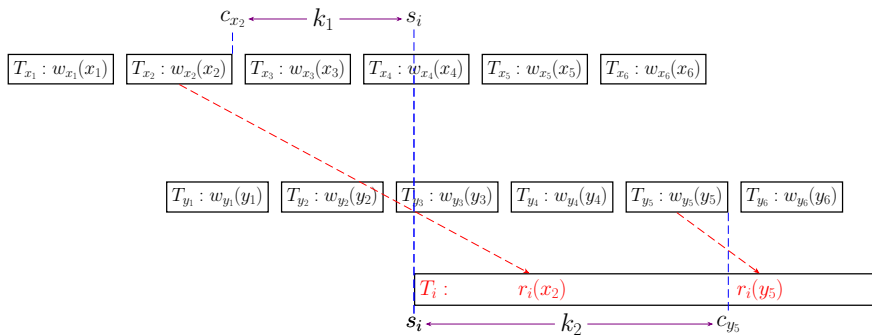
$$\boxed{T_{x_1} : w_{x_1}(x_1)} \quad \boxed{T_{x_2} : w_{x_2}(x_2)} \quad \boxed{T_{x_3} : w_{x_3}(x_3)} \quad \boxed{T_{x_4} : w_{x_4}(x_4)} \quad \boxed{T_{x_5} : w_{x_5}(x_5)} \quad \boxed{T_{x_6} : w_{x_6}(x_6)}$$


$$\boxed{T_{y_1} : w_{y_1}(y_1)} \quad \boxed{T_{y_2} : w_{y_2}(y_2)} \quad \boxed{T_{y_3} : w_{y_3}(y_3)} \quad \boxed{T_{y_4} : w_{y_4}(y_4)} \quad \boxed{T_{y_5} : w_{y_5}(y_5)} \quad \boxed{T_{y_6} : w_{y_6}(y_6)}$$

$$\boxed{T_i : \quad r_i(x_2) \quad r_i(y_5)}$$









$(k_1\text{-BV})$ 

$$\forall r_i(x_j), w_k(x_k), c_k \in h : \left( c_j \in h \wedge \bigwedge_{k=1}^m (c_j \prec_h c_k \prec_h s_i) \right) \Rightarrow m < k.$$

 $(k_2\text{-FV})$ 

$$\forall r_i(x_j), w_k(x_k), c_k \in h : \left( c_j \in h \wedge \bigwedge_{k=1}^m (s_i \prec_h c_k \prec_h c_j) \right) \Rightarrow m \leq k.$$

 $(k_3\text{-SV})$ 

$$\forall r_i(x_j), r_i(y_l), w_k(x_k), c_k \in h : \left( \bigwedge_{k=1}^m (c_j \prec_h c_k \prec_h c_l) \right) \Rightarrow m \leq k_3.$$

$$h \in \text{RVSI} \iff h \in k_1\text{-BV} \cap k_2\text{-FV} \cap k_3\text{-SV} \cap \text{WCF}$$

$$\text{RVSI}(\infty, \infty, \infty) = \text{RC} \quad \text{RVSI}(1, 0, *) = \text{SI}$$

# Parameterized and Runtime-tunable Snapshot Isolation

## RVSI: Relaxed Version Snapshot Isolation

- 1 Motivation for RVSI
- 2 Definition of RVSI
- 3 CHAMELEON Prototype and RVSI Protocol
- 4 Experimental Evaluation

# CHAMELEON Protocol

CHAMELEON:

A prototype **partitioned replicated**  
distributed transactional **key-value** store


# CHAMELEON Prototype

Classic **key-value** data model


Key: (row key, column key)



# CHAMELEON Prototype




# CHAMELEON Prototype




Keys are **partitioned** within a single datacenter.

# CHAMELEON Prototype




Each key is **replicated** in a master-slave manner across datacenters.

# CHAMELEON Prototype

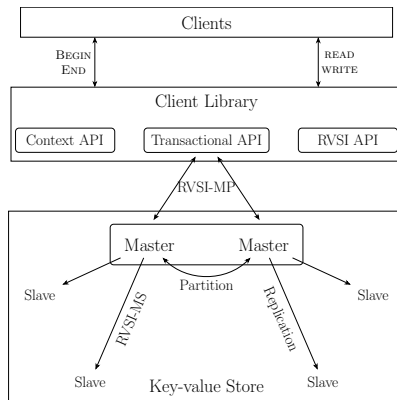


Transactions are first executed and committed on the masters, and are then asynchronously propagated to slaves.

# CHAMELEON Protocol

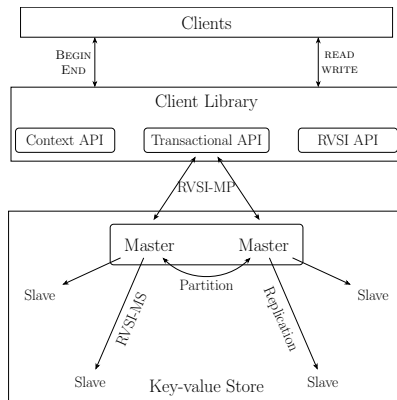


# CHAMELEON Protocol



Partitioned replicated transactional key-value store

# CHAMELEON Protocol



Client library

# CHAMELEON Protocol

```
// Initialize keys (ck, ck1, and ck2) here
ITx tx = new RVSI Tx(** context **);


tx.begin();

// Read and write
ITsCell tsCell = tx.read(ck);
ITsCell tsCell1 = tx.read(ck1);
tx.write(ck1, new Cell("R1C1"));
ITsCell tsCell2 = tx.read(ck2);

// Specify RVSI specs. (e.g., SVSpec)
RVSI Spec sv = new SVSpec();
sv.addSpec({ck, ck1, ck2}, 2);
tx.collectRVSI Spec(sv);
```



# CHAMELEON Protocol



RVSI protocol: RVSI-MS + RVSI-MP

# RVSI-MS Protocol

**RVSI-MS:** RVSI protocol for master-slave replication

In terms of *event* generation and handling:

**Clients:** BEGIN, READ, WRITE, END

**Master:** START, COMMIT, SEND

**Slaves:** RECEIVE

# RVSI-MS Protocol

TikZ overlay for RVSI-MS

# RVSI-MS Protocol

Calculating version constraints:

# RVSI-MP Protocol

Distributed transactions spanning multiple masters need to be committed atomically.

# RVSI-MP Protocol

Distributed transactions spanning multiple masters need to be committed atomically.  
Using the two-phase commit (2PC) protocol.

# RVSI-MP Protocol

Assumes a timestamp oracle:

Clients:

Masters:

# RVSI-MP Protocol

RVSI version constraints in 2PC protocol:

$k_1$ -BV:

$k_2$ -FV:

$k_3$ -SV:



# RVSI-MS Protocol

---

**Algorithm 1** RVSI-MS: RVSI Protocol for Replication (for Executing Transaction  $T$ ).

---

***Client-side methods:***

- 1: **procedure** BEGIN()
- 2:      $T.sts \leftarrow \mathbf{rpc-call}$  START() at master  $\mathcal{M}$
- 3: **procedure** READ( $x$ )
- 4:      $x.ver \leftarrow \mathbf{rpc-call}$  READ( $x$ ) at any site
- 5: **procedure** WRITE( $x, v$ )
- 6:     add  $(x, v)$  to  $T.writes$
- 7: **procedure** END( $T$ )
- 8:      $T.vc \leftarrow \mathbf{ADD-VC}()$
- 9:      $c/a \leftarrow \mathbf{rpc-call}$  COMMIT( $T.writes, T.vc$ ) at  $\mathcal{M}$

# RVSI-MP Protocol

---

**Algorithm 2** RVSI-MP: RVSI Protocol for Partition (for Executing Transaction  $T$ ).

---

***Client-side methods:***

- 1: **procedure** BEGIN()
- 2:     **return** **rpc-call** GETTS() at  $\mathcal{T}$
- 3: **procedure** END()
- 4:      $T.vc \leftarrow \text{ADD-VC}()$
- 5:      $c/a \leftarrow \text{rpc-call C-COMMIT}(T.writes, T.vc)$  at  $\mathcal{C}$

---

***Timestamp oracle methods:***

$\mathcal{T}.ts$ : for start-timestamps and commit-timestamps

- 6: **procedure** GETTS()
- 7:     **return**  $++\mathcal{T}.ts$



# Parameterized and Runtime-tunable Snapshot Isolation

## RVSI: Relaxed Version Snapshot Isolation

- 1 Motivation for RVSI
- 2 Definition of RVSI
- 3 CHAMELEON Prototype and RVSI Protocol
- 4 Experimental Evaluation

# Impacts of RVSI specification on the *transaction abort rates* in various scenarios

## Impacts of RVSI specification on the *transaction abort rates* in various scenarios

### Performance?

- ▶ Not done yet in this work
- ▶ CHAMELEON prototype is ...

Transactions abort for two reasons:

- ▶ “vc-aborted”: RVSI version constraints violated
- ▶ “wcf-aborted”: the WCF property violated

Transactions abort for two reasons:

- ▶ “vc-aborted”: RVSI version constraints violated
- ▶ “wcf-aborted”: the WCF property violated

Transaction abort rates due to “vc-aborted” are *sensitive* to different values of  $k_1$ ,  $k_2$ , or  $k_3$ ,



Transactions abort for two reasons:

- ▶ “vc-aborted”: RVSI version constraints violated
- ▶ “wcf-aborted”: the WCF property violated

Transaction abort rates due to “vc-aborted” are *sensitive* to different values of  $k_1$ ,  $k_2$ , or  $k_3$ , but those due to “wcf-aborted” are not.

Transactions abort for two reasons:

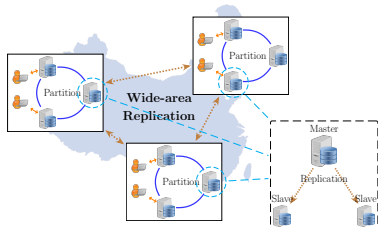
- ▶ “vc-aborted”: RVSI version constraints violated
- ▶ “wcf-aborted”: the WCF property violated

Transaction abort rates due to “vc-aborted” are *sensitive* to different values of  $k_1$ ,  $k_2$ , or  $k_3$ , but those due to “wcf-aborted” are not.

$$h \in \text{RVSI} \iff h \in k_1\text{-BV} \cap k_2\text{-FV} \cap k_3\text{-SV} \cap \text{WCF}.$$

## CHAMELEON prototype on Aliyun:

- ▶ 3 datacenters <sup>1</sup>
- ▶ 3 nodes in each datacenter
- ▶ Partition & Replication
- ▶ Clients in our lab <sup>2</sup>



<sup>1</sup> Located in East China, North China, and South China, respectively.

<sup>2</sup> Located in East China.

<sup>3</sup> <https://github.com/hengxin/aliyun-ping-traces>

## CHAMELEON prototype on Aliyun:

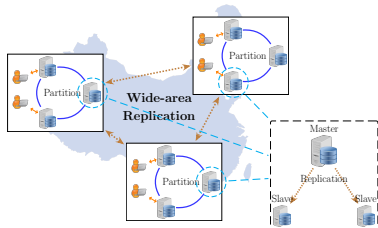
- ▶ 3 datacenters <sup>1</sup>
- ▶ 3 nodes in each datacenter
- ▶ Partition & Replication
- ▶ Clients in our lab <sup>2</sup>

(One-way) delays among nodes <sup>3</sup>:

Within datacenter: 1 ~ 2ms

Across datacenters: 15 ~ 25ms

Clients to nodes: 15 ~ 20ms



<sup>1</sup> Located in East China, North China, and South China, respectively.


<sup>2</sup> Located in East China.


<sup>3</sup> <https://github.com/hengxin/aliyun-ping-traces>

Table: Three categories of workload parameters for experiments on Aliyun.

Parameter		Value	
Transaction-related	#keys	$25 = 5 \text{ (rows)} \times 5 \text{ (columns)}$	
	#clients	5, 10, 15, 20, 25, 30	
	#txs/client	1000	
	#ops/tx	$\sim \text{Binomial}(20, 0.5)$	n
	rwRatio	1:2, 1:1, 4:1	
	zipfExponent	1	
Execution-related	minInterval	0ms	
	maxInterval	10ms	
	meanInterval	5ms	
RVSI-related	$(k_1, k_2, k_3)$	$(1,0,0) (1,1,0) (1,1,1)$ $(2,0,0) (2,0,1) (2,1,1)$	


under read-frequent workloads






The transaction abort rates due to “vc-aborted”






The transaction abort rates due to “vc-aborted” can be **greatly reduced** by **slightly** increasing the values of  $k_1$ ,  $k_2$ , or  $k_3$ :

$$vc(1, 0, 0) = 0.1994 \implies vc(2, 1, 1) = 0.0091 \quad (\#clients = 30)$$





Most “vc-aborted” transactions abort because of violating  $k_2\text{-FV}$ .

$$fv(1, 0, 0) = 0.1889 \implies fv(2, 0, 0) = 0.1866 \implies fv(1, \mathbf{1}, 0) = 0.0064$$

Question: when does  $k_1$  for  $k_1$ -BV take effect?

It seems that  $k_1$ -BV has *little* impact on the transaction abort rates.

Question: when does  $k_1$  for  $k_1$ -BV take effect?


It seems that  $k_1$ -BV has *little* impact on the transaction abort rates.

It may be the case in the Aliyun scenarios.

What about other scenarios?

Three types of delays for **controlled experiments** on local hosts.

Types	Values (ms)	Explanation
<b>issueDelay</b>	5, 8, 10, 12, 15, 20	delays between clients and replicas
<b>replDelay</b>	5, 10, 15, 20, 30	delays between masters and slaves
<b>2pcDelay</b>	10, 20, 30, 40, 50	delays among masters



When the “**issueDelay**” gets shorter,  
the impacts of  $k_2$ -FV go weaker,  
and the impacts of  $k_1$ -BV have begun to emerge.

issueDelay = 20ms:  $bv(1, 0, 0) = 0.0057$      $fv(1, 0, 0) = 0.0251$

issueDelay = 15ms:  $bv(1, 0, 0) = 0.08225$      $fv(1, 0, 0) = 0.0393$

issueDelay = 5ms:  $bv(1, 0, 0) = 0.1716$      $fv(1, 0, 0) = 0.0045$



issueDelay = 20ms:  $bv(1, 0, 0) = 0.0057$      $fv(1, 0, 0) = 0.0251$

issueDelay = 15ms:  $bv(1, 0, 0) = 0.08225$      $fv(1, 0, 0) = 0.0393$

issueDelay = 5ms:  $bv(1, 0, 0) = 0.1716$      $fv(1, 0, 0) = 0.0045$

larger issueDelay  $\implies$  longer transaction

more concurrent transactions

more likely to obtain data versions updated by concurrent transactions

more sensitive to  $k_2$ -FV

Generally, RVSI **helps to reduce** the transaction abort rates when applications are willing to tolerate certain anomalies.

Generally, RVSI **helps to reduce** the transaction abort rates when applications are willing to tolerate certain anomalies.

$k_2$ -FV: In the Aliyun scenarios, most transactions have been aborted because of violating  $k_2$ -FV.

Generally, RVSI **helps to reduce** the transaction abort rates when applications are willing to tolerate certain anomalies.

$k_2$ -FV: In the Aliyun scenarios, most transactions have been aborted because of violating  $k_2$ -FV.

$k_1$ -BV: In controlled experiments, the impacts of  $k_1$ -BV emerge when the issueDelay gets shorter.

Generally, RVSI **helps to reduce** the transaction abort rates when applications are willing to tolerate certain anomalies.

$k_2$ -FV: In the Aliyun scenarios, most transactions have been aborted because of violating  $k_2$ -FV.

$k_1$ -BV: In controlled experiments, the impacts of  $k_1$ -BV emerge when the issueDelay gets shorter.

$k_3$ -SV: Complex and challenging (involving multiple data items)